

Gate 6B - Review

Open only after saved fresh retest | RC1

How to use this review

Open only after the attempt is saved. Diagnose the failed competency, repair that competency in a copy of the attempt, and preserve the original evidence. This review intentionally does not contain a finished architecture or code solution.

Review dimensions

- Risk identification and priority: correctness, security, reliability, cost and operability
- Fabric/OneLake/Lakehouse architecture and medallion responsibilities
- Incremental load, watermark, merge and rerun/idempotency safety
- Automated data-quality gate and PII masking behavior
- Identity, secret handling, RBAC and least privilege
- CI/CD, environment promotion and rollback reasoning
- Monitoring, alerting, lineage, incident response and runbook quality
- Terraform/IaC foundation and configuration hygiene
- Cost/scaling trade-offs and technical defense

Repair workflow

1. Map every failed item to ARCHITECTURE, DATA, SECURITY, CI/CD, OBSERVABILITY, IAC or EXPLAIN.
2. State the root cause in one sentence and the production consequence in one sentence.
3. Repair only the weak competency in a copy; do not overwrite the original attempt.
4. Rerun the checker and save the new output and changed-file list.
5. Defend the repair: why it is safer, how rollback/recovery works, and what changes at 10x scale.

Retest rule

If a critical floor still fails, repair the underlying competency and complete a new unseen exercise before another Gate decision. Do not rehearse this exact domain.

A high overall score cannot compensate for a failed critical floor such as leaked secrets, unsafe reruns, incorrect data, missing recovery, or unjustified broad access.