

M11 Protected Review & Repair

LOCKED: open a lesson section only after a genuine saved independent attempt. This is diagnosis/repair, not a substitute first attempt.

DE11-01 - Review after attempt

Expected proof: Your explanation must not call Parquet a database. `orders.csv` has 12 rows, quantity 24, amount 2780.00, four regions and four dates.

Common repair: If row-vs-columnar still feels abstract, draw six columns and four rows on paper, then circle the values needed for `SUM(amount) GROUP BY region`. Count how many cells are relevant.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-02 - Review after attempt

Expected proof: For 12 rows written with a row-group target of 4, expect 3 row groups in this small controlled file. Do not generalize this exact count to production writers.

Common repair: If PyArrow import fails, verify the active interpreter and run `python -m pip install "pyarrow==25.0.*"; use the PAR1 probe meanwhile. Do not install Spark to solve a PyArrow problem.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-03 - Review after attempt

Expected proof: For amount ≥ 250 on the course data, qualifying orders are O1002, O1004, O1008, O1010, O1012; count 5 and amount 1610.00.

Common repair: If a filtered read returns unexpected rows, prove the predicate against CSV first. Separate a logic error from a storage/engine optimization issue.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-04 - Review after attempt

Expected proof: The classroom layout has 12 tiny files and 1 compacted file while preserving the 12-row business dataset. Normal amount remains 2780.00.

Common repair: If you think "one huge file is always best," list concurrency, parallelism and rewrite-cost trade-offs. The goal is healthy file sizing, not minimum file count.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-05 - Review after attempt

Expected proof: Course data distinct counts: region=4, order_date=4, order_id=12. A partition-by-order_id plan creates the clearest over-partitioning warning here.

Common repair: If you cannot choose a partition key, start from the most common selective query boundaries and data volume. Never choose only because a column "looks categorical."

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-06 - Review after attempt

Expected proof: Your answer must name trade-offs. "Lakehouse is newer so always better" is an automatic reasoning defect.

Common repair: If the terms blur together, separate physical storage (files/object store) from table-management/catalog/query layers.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-07 - Review after attempt

Expected proof: Simulator versions are [0,1,2]. V0 has 4 rows / spend 820.0; V1 has 5 rows / spend 1050.0; V2 adds `segment` without changing row count.

Common repair: If you equate Delta with "Parquet plus a JSON folder," focus on the protocol: correctness comes from committed log semantics, not merely the file names.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-08 - Review after attempt

Expected proof: Expected simulator MERGE: updated=[C2], inserted=[C5], rows=5, spend=1050.0. Multiple source records for one target key must be resolved before claiming safe current-state logic.

Common repair: If target rows unexpectedly multiply, stop and test key uniqueness on both target and intended source update grain before changing the MERGE clause.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-09 - Review after attempt

Expected proof: V2 schema is [id,name,tier,spend,segment] and v2_rows remains 5. Row-count stability alone is not enough; verify schema and key/control totals too.

Common repair: If "schema evolution" sounds like permission to ignore drift, rewrite the rule as: detect -> classify -> approve -> migrate -> validate -> monitor.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.

DE11-10 - Review after attempt

Expected proof: Time-travel control: v0 C2 tier=Gold; current v2 C2 tier=Platinum. VACUUM remains dry-run only. Your medallion checklist must preserve raw evidence and include row/key/control reconciliation.

Common repair: If time travel and backups seem identical, separate table-version history from an independent backup/disaster-recovery copy. They solve overlapping but different failure cases.

Targeted changed micro-retry: Rewrite your answer or rerun the check with a changed predicate/key/partition choice/version. Do not copy the original task output. Then close this Review.

Pass evidence: independent result + validation + explain-back must all be coherent. A numeric percentage never infers mastery.