

M15 PROTECTED REVIEW + REPAIR

LOCKED UNTIL A GENUINE SAVED ATTEMPT EXISTS. Diagnose the first weak boundary; do not use this as a pre-attempt answer sheet.

DE15-01 — Review

What strong evidence shows

Expected proof: one architecture map that labels source, persistent state, compute/orchestration, consumer, identity boundary and monitoring signal.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design a changed Fabric architecture for a small parcel company with API + files + analytics consumers. Name state, compute, security and monitoring boundaries.

DE15-02 — Review

What strong evidence shows

Expected proof: a permission matrix with at least one deliberately refused grant and one test performed as the intended narrow identity.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Create a least-privilege permission matrix for a Viewer analyst, Contributor engineer and service identity; defend every grant.

DE15-03 — Review

What strong evidence shows

Expected proof: three copy-vs-shortcut decisions with freshness, ownership and replay/recovery reasons.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Choose copy or shortcut for a changed source set and record one failure/recovery concern per choice.

DE15-04 — Review

What strong evidence shows

Expected proof: a workspace/item inventory naming purpose, owner, input/output and validation evidence for every item.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Build a fresh workspace/item inventory from a blank template without looking at the guided example.

DE15-05 — Review

What strong evidence shows

Expected proof: explicit Files, Tables and shortcut paths plus a sentence explaining which bytes are copied and which are referenced.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design fresh OneLake Files/Tables/shortcut paths and explain how SQL/Spark would see each.

DE15-06 — Review

What strong evidence shows

Expected proof: table grain + schema contract + bad-row rules + 7 raw / 4 accepted / 3 quarantined / 3300 control for the normal fixture.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Define a fresh lakehouse Delta table contract and validation for a changed dataset.

DE15-07 — Review

What strong evidence shows

Expected proof: rerunnable notebook flow with parameters/config separated from transformation logic and a clean-start validation plan.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design a rerunnable notebook flow from a blank file: parameters → read → validate → transform → write → reconcile.

DE15-08 — Review

What strong evidence shows

Expected proof: compatibility checklist that explicitly records local Spark 4.2/Delta 4.4 versus Fabric Runtime 2.0 Spark 4.1/Delta 4.2.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Write a runtime-compatibility checklist for moving local Spark/Delta code into Fabric Runtime 2.0.

DE15-09 — Review

What strong evidence shows

Expected proof: dependency graph, parameter, safe-retry decision, blocking failure route and checkpoint/publish rule.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design an independent pipeline with at least four activities, dependency, parameter, retry and failure branch.

DE15-10 — Review

What strong evidence shows

Expected proof: two read-only validation queries plus a justified lakehouse-vs-Warehouse choice for changed workloads.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Choose lakehouse SQL endpoint vs Warehouse for three changed workloads and defend each choice.

DE15-11 — Review

What strong evidence shows

Expected proof: Bronze/Silver/Gold contracts with layer grain, acceptance/rejection behavior, lineage and reconciliation.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Define Bronze/Silver/Gold contracts for a changed domain with explicit grains and rejected-row handling.

DE15-12 — Review

What strong evidence shows

Expected proof: base controls, update/insert controls, 5 rows / 4250 after merge, exact-rerun no-op and failed-publish checkpoint rule.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design a changed incremental load where a failed publish must not advance progress. State the atomic boundary.

DE15-13 — Review

What strong evidence shows

Expected proof: identity -> permission -> resource matrix, no literal secret, and evidence for one allowed plus one denied access path.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Find and repair three security mistakes in the changed plan: broad role, hard-coded secret and unclear external connection identity.

DE15-14 — Review

What strong evidence shows

Expected proof: monitoring record covering run status, freshness, quality/reconciliation, duration, failure detail and capacity/cost signal.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Create a monitoring checklist covering freshness, quality, failure, duration and capacity/cost signals.

DE15-15 — Review

What strong evidence shows

Expected proof: dev/test/prod mapping, Git branch/review path, environment configuration separation, validation-after-deploy and rollback plan.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Design a Git + deployment pipeline flow with dev/test/prod configuration separation and rollback evidence.

DE15-16 — Review

What strong evidence shows

Expected proof: end-to-end architecture/evidence index that distinguishes actual Fabric observations from offline design/simulator evidence.

Diagnostic questions

1. Did you name the actual Fabric/cloud object or clearly label the work OFFLINE PLAN? 2. Did you distinguish identity/permission from data/state? 3. Did you validate an observable count/status/path/permission/version rather than trust a green badge? 4. Did you preserve the first attempt and runtime limitation?

Targeted repair

Repair only the first missing/incorrect boundary. If the problem is a platform/runtime issue, record the exact symptom and current official reference rather than inventing success. Then close Review.

Assigned fresh retry

Produce the complete clinic retest architecture/evidence plan from a blank template and explain it without the book open.