

M18 / ASSESSMENT

M18 Module Revision Test

Standalone Zero-to-Job-Ready Data Engineer 2026

Study, practise, check and explain.

A complete learning route

Start with the lesson video. Read the worked explanation beside it, pause to reproduce the example, then solve the supplied task. The books contain your study notes; keep your own work and results as evidence.

Use the protected answers after saving your first attempt. A changed retry shows what you can now do independently. There is no forced timer during learning.

One course. M00-M19. 264 lessons and 50 project phases.

Contents

Tap a topic or use the PDF bookmarks. These numbers are the actual pages in this edition.

Read this first 3

Read this first

Rule: Untimed mastery first. Close the portfolio manual, repositories and protected answers. Save your genuine attempt before review.

1. Explain why a portfolio pipeline that produces the correct final CSV once can still fail the M18 standard.
2. Design an immutable raw-landing identity and explain changed-content collision behavior.
3. A source batch has 1,000 rows, 12 invalid business records and a sink failure after 700 writes. Separate quarantine from infrastructure recovery and state ownership.
4. Show three source-to-target reconciliation checks for an order pipeline.
5. A fact join doubles revenue. Diagnose the likely grain/cardinality problem and prove the repair.
6. Explain why an incremental watermark using `max(updated_at)` can still miss data, and give a safer bounded strategy.
7. Define a dbt-style staging/fact/dimension structure and where tests belong.
8. Describe an orchestration failure where progress outruns durable output; give the safe repair.
9. For Spark, explain how partition count, shuffle and skew evidence change a performance decision.
10. Define Bronze/Silver/Gold as contracts rather than folder names.
11. An additive schema field appears. Describe compatibility, quarantine/versioning and downstream validation.
12. Map a local lakehouse project to Fabric items without claiming a tenant run you did not execute.
13. Choose CI checks that block deployment versus observability checks that alert after deployment.
14. Write the minimum contents of a data-pipeline runbook.
15. Explain what a clean-clone proof establishes and what it does not establish.
16. Give a two-minute technical defense structure for a portfolio project.
17. Capstone incident: duplicated payment event plus stale inventory CDC change. Explain safe diagnosis, replay and uniqueness/version controls.
18. State when Spark/Delta is justified in the capstone and when it is unnecessary complexity.
19. Explain how least privilege and secret boundaries should appear in a Fabric/cloud deployment design.
20. List the evidence bundle you would show an interviewer to prove this is your own engineering work.