

P1 - Reliable Batch Pipeline - Review

Open after a complete saved attempt | RC2

Review dimensions

- Correctness and grain
- Reproducibility and rerun safety
- Data quality and reconciliation
- Testing and failure handling
- Security/config hygiene
- Operational recovery/backfill
- Documentation and architecture
- Technical explanation

Repair rule

Repair the failed competency in a copy of your submitted project. Preserve the first failing evidence. Do not replace the whole project with a new answer.

Defense prompts

- What is the grain of the trusted output?
- What breaks if this pipeline runs twice?
- Which failure is detected earliest and how?
- How do you recover from a partial run?
- What would you change at 10x data volume?
- Which design choice would you revisit first in production?